

# R on the HPC and parallel computing

Eleanor Zhang



Date: 08/18/2022

# Aims



Part I: Parallel computing in R: *parallel* package



Part II: Parallel computing using batch jobs on HPC cluster



Part III: Practice on your own

# How does R interact with operational system?

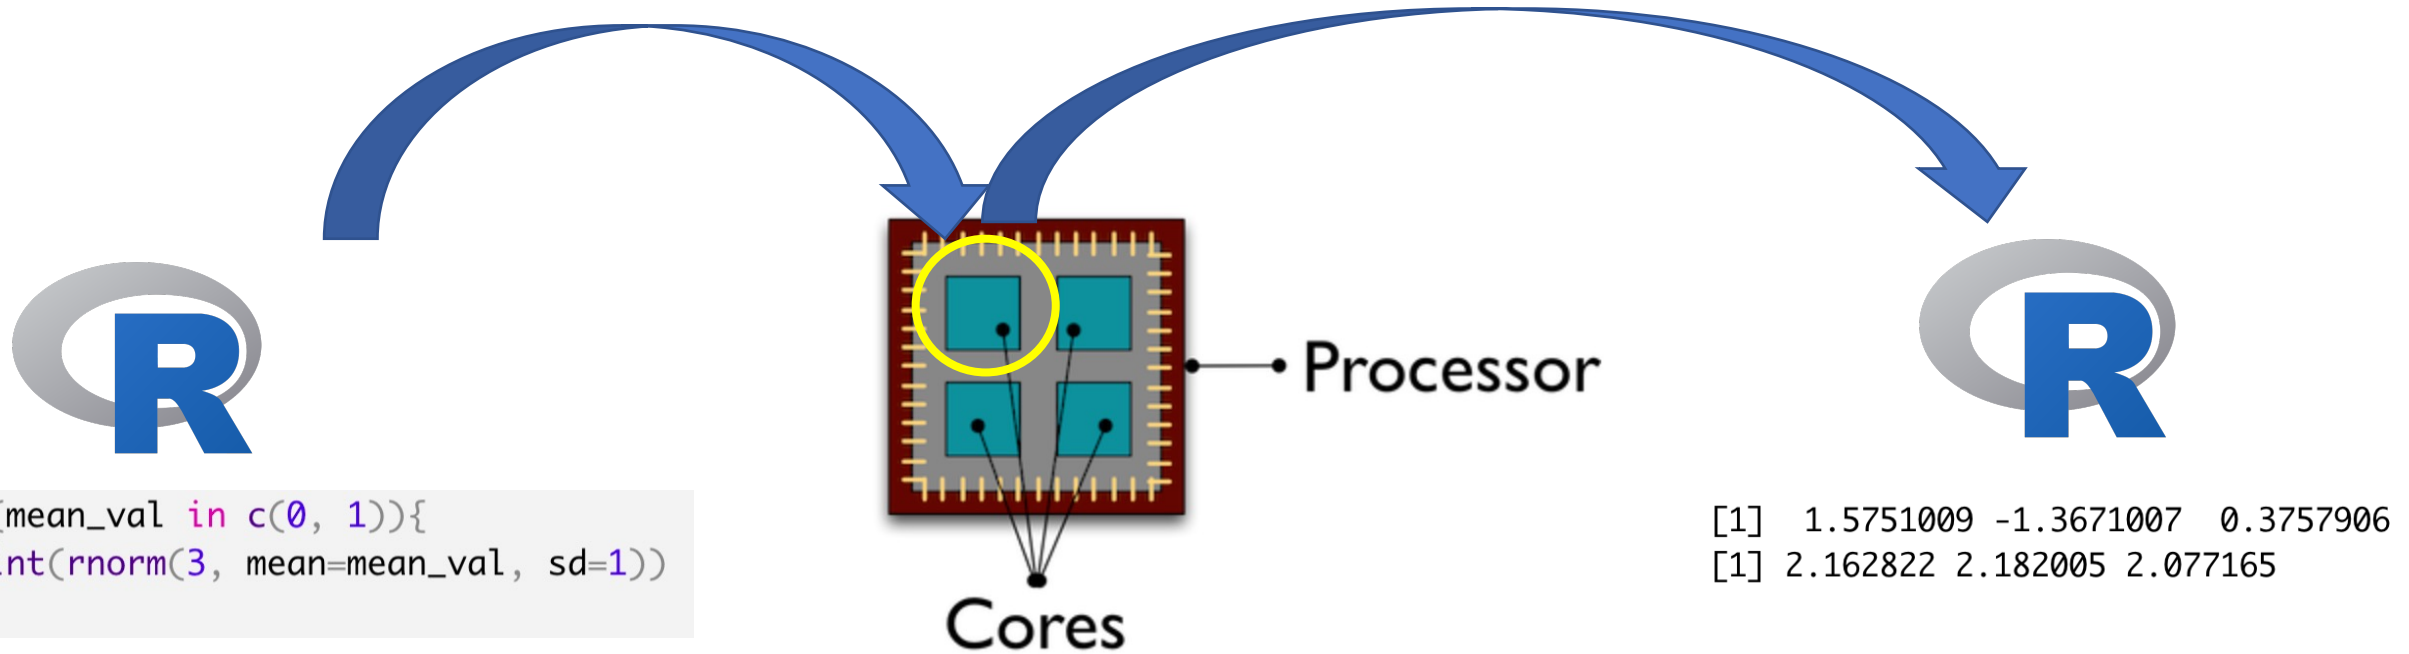


```
for (mean_val in c(0, 1)){  
  print(rnorm(3, mean=mean_val, sd=1))  
}
```

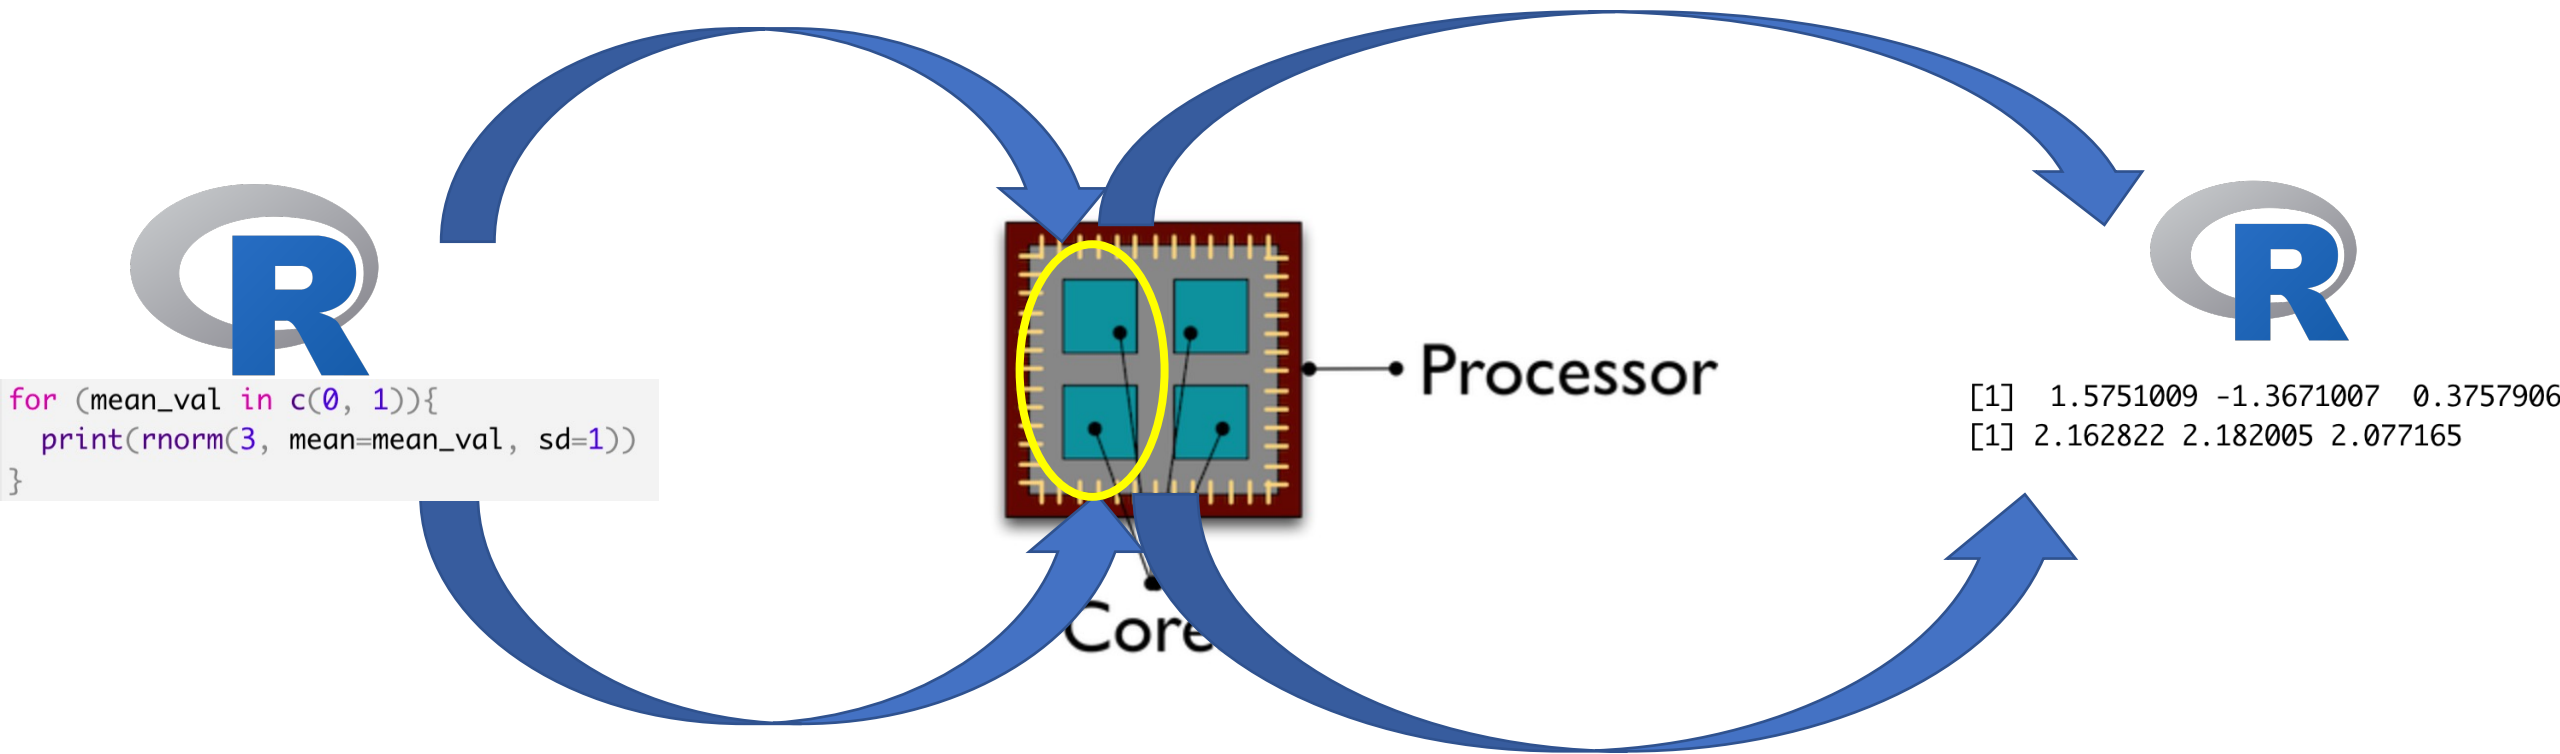
- Central processing unit (CPU): the “brain”



# R uses single core by default



# R in parallel using multi-core for independent tasks



# *Parallel* package:

Two functions to implement multi-cores for independent tasks:

	Forking method	Socket method
R function	mclapply()	parLapply()

→ Go to parallel.Rmd

# *Parallel* package:

Two functions to implement multi-cores for independent tasks:

	Forking method	Socket method
R function	mcapply()	parLapply()
Pros and Cons	Simple; Not applicable to Windows users	Structured code; Compatible with any system



# Overhead in parallel computing



- Amdahl's law: speed up in computing time is NOT proportional to number of cores
- Some computation tasks not worth parallelization
- Try out and see what number of cores to use

# Happy with your personal computer?



# Hardware on personal computer


- CPU and RAM:

Processor 2.3 GHz Dual-Core Intel Core i5  
Memory 8 GB 2133 MHz LPDDR3

- By default, R uses single core and single thread
- R stores and manipulate all objects in RAM in R session

# Hardware on personal computer


- Storage on disk



- R write to disk: `save()`, `write_tsv()`, etc.

# High performance computing (HPC)

- HPC aggregates the resources from individual computers (known as *nodes*) into a cluster that works together to perform advanced, specialized computing jobs.




# HPC at USC

CARC: Center for Advanced Research Computing

- <https://www.carc.usc.edu/>
- <https://www.carc.usc.edu/user-information/user-guides/hpc-basics/discovery-resources>
- Provides workshop and weekly office hours

# CARC computing resource

- Login node (Head node)
- Compute node wired together (~400)




# CARC computing resource

- Login node (Head node)
- Compute node wired together (~400)

One rack in HPC cluster




Multiple racks in HPC cluster





# CARC computing resource

- Login node (Head node)
- Compute node wired together (~400)
- Networked storage ( ~12 PB)
  - Disk array and file servers
- Slurm job scheduler
- **Discovery** and Endeavour



# Discovery account

- PI can add a member to a project
- For this R bootcamp session, Eric added us to a project and we should have access to `/project/ekawaguc_906`
- Username: USC NetID (eg. `zzhang39@usc.edu`)

# Connecting to Discovery login node

Method 1: Secure shell command line interface

- On Mac: Terminal
- On Windows: OpenSSH, Putty, etc

```
ssh <username>@discovery.usc.edu
```

Method 2: Use OnDemand interface

<https://www.carc.usc.edu/user-information/user-guides/hpc-basics/getting-started-ondemand>

Requirement:

- Network connection
- VPN if off-campus
- DUO Two Factor Authentication: see <https://itservices.usc.edu/duo/> to set up


# Navigate file system

myquota

- Home directory (private):  
/home1/<username>
- Two scratch directories:  
/scratch1/<username>  
/scratch2/<username>
- Project: /project/<pi\_username\_id>

# Slurm: the “manager”

- Slurm is a cluster management system that can allocate and monitor computing resources




# Request for interactive compute node

In shell:

```
salloc --time=1:00:00 --mem=8Gb
```

*Please do not run heavy task on login node!*



# Using R on HPC

## Method 1: Pre-built R module by CARC

```
module purge                # remove all currently loaded modules
module load gcc/11.3.0      # load compiler
module load openblas/0.3.20 # load linear algebra subprogram
module load r/4.2.1
```

Method 2: Install and manage R using other distributions: eg. Anaconda (Practice in the end)

# Run R script on command line


- Create R script (with or without arguments)
  - Text editor: vim, Emacs, nano, etc
- Run command in shell:

```
Rscript <scriptname> arg1 arg2 ...
```



# Submit batch jobs using Slurm

- Batch job: commands that will be executed with little interaction between user and system
- Use job script to submit batch jobs



```
#!/bin/bash
#SBATCH --ntasks=1
#SBATCH --mem-per-cpu=1G
#SBATCH --time=01:00:00

#set path, environment variables
module load gcc/8.3.0
module load python/3.7.4

#run program
python3 my_cool_script.py
```

# Example of using batch jobs

- Previous example

Calculate bootstrap standard errors for slope in logistic regression:

Species ~ Sepal.Length

- Now let's calculate bootstrap standard errors of slope for each feature variable:

Species ~ Sepal.Width

Species ~ Sepal.Length

Species ~ Petal.Width

Species ~ Sepal.Length

# Practice

- Set up DUO two factor: <https://itservices.usc.edu/duo/>
- Go to practice\_instruction.html file and follow the steps to practice on your own
- Let me know if you have questions

# Reference

R parallel computing:

<https://nceas.github.io/oss-lessons/parallel-computing-in-r/parallel-computing-in-r.html>

USC HPC workshops:

- Discovery cluster overview:

<https://www.youtube.com/watch?v=gtX6VM2UkTo>

- Using R: <https://www.youtube.com/watch?v=3sRy2Lx8Hlc&t=658s>